

Make the next step obvious.

The product direction is strong. The experience is not yet consistently easy, coherent, or complete enough to deliver that promise to every user.

Keep the interactive funnel, contextual Kai, paper practice, community, and the Look / Decide / Take trade flow. Focus the next build on continuity, a small family of distinctive visual components, and clear guidance matched to each person.

The experience to build

I understand what I am looking at. I know what I can do next. Kai helps me make sense of it. I can see what changed, and my progress is still here when I return.

Priority	Most consequential findings
P0 - complete first	Only 1 of 31 training lessons is playable; training progress is device-wide; Home conversation selection is not connected to conversation identity.
P1 - next release	Core trade levels are collapsed; community navigation changes global mode; beginner Home starts with market prose; accessibility settings are not consumed by the UI.
P2 - refine	Unify statuses and typography, simplify Account, improve learning milestones, and make quiet / empty / error states useful.

What this audit establishes

This is a source-backed product audit with visual inspection. It identifies implementation defects, UX friction, and design recommendations separately. It does not claim measured conversion improvement or a completed authenticated device test.

No production code, account data, orders, or community messages were changed during this audit.

What was reviewed

Source baseline: kwayclawdbot/cheatcode-ai-app, origin/feat/everything-on-57, commit 7f7580196a351e14415708298ce6f8c8f85c1e9d. This is the newer feature branch fetched for the audit, not the older origin/main.

Reviewed the major customer journeys across the mobile source, relevant API rules, the current site funnel, and 14 repository proof images. The mobile route inventory contains 63 non-layout route files, including 54 customer-route candidates after excluding admin and development harnesses. Inventory is not a claim that every route was runtime-tested.

Evidence	What it can establish	Boundary
Source	Rendered components, routing, state ownership, API handoffs and available content.	Some defects are source-confirmed but still require reproduction in a signed-in release build.
Repository images	Layout, visual hierarchy and historical design intent.	Proof images span earlier commits. Current source takes precedence where they disagree.
Live public web	Welcome and signup at cheatcode-ai-app.vercel.app on 8 September.	No signed-in session was available. Deployment commit is unknown; native behavior was not tested.
Inspection / calculation	Token contrast and authored versus playable lesson counts.	No full accessibility certification, performance benchmark, cohort analysis or conversion experiment.

A significant delivery distinction

The reusable component kit we made is in the local feat/trade-ui-kit work and the previously delivered kit. Its TradeMap, RiskRewardRuler, SetupPreview and related components are not integrated into the current audited remote app screens. A component demo is not evidence of app-wide adoption.

The current Trade route does use TradePortalV2. Earlier screenshots showing Kai / Alert / Plan / Community tabs or a microphone composer are historical; this report does not count those removed designs as current defects. [S4, S13]

Build one product for different needs

Personalization should change the first useful object, explanation depth and suggested action. It should not make users relearn navigation or equate trading frequency with proven skill.

Member	First useful view	Primary action	Kai behavior
Brand new	One tactile ownership lesson, with a short reason to try it.	Try the lesson	Explain one term at a time; invite questions.
Long-term investor	One company explained through business, horizon and risks.	Understand this company	Teach investing without forcing intraday terminology.
Developing	Next skill plus one simple, fully explained setup.	Practise the plan	Ask the member to identify entry and invalidation.
Trade Ready	Actionable setup or position needing attention, with its state.	Review setup / position	Levels first; challenge reasoning on request.
Returning member	Resume unfinished work; explain what changed since last visit.	Continue	Restore exact context and avoid repeating onboarding.

Separate the four concepts

Goal: Invest / Swing / Day trade. Guidance: Guided / Concise, adjustable any time. Readiness: demonstrated learning milestones, with self-reported placement identified. Belt: community track record, with scoring explained. Free / paid access is a separate entitlement dimension.

Today there are start_answer, experience, explanation_level, stage and belts, with partially overlapping labels. An active trader can be placed Trade Ready while the draft experience remains New unless they change it later. Account also offers Beginner / Intermediate / Advanced / Family for explanations. Reconcile these inputs and make the outcome visible. [S2, S3, S8]

Give guidance a visible control near Kai. Keep the investor path available at every readiness stage. Do not hide basic explanations from experienced members or push beginners toward complex instruments.

Onboarding should deliver a first win

F01 - P1 / Repeated questions and a broken promise of continuity

Evidence. The app asks placement, goal, risk, experience/focus, username and a Kai plan before Home. Placement and experience are separate; the draft is in React state. Site choices reach an email draft, not the app profile. [S2, S10]

Why it matters. A visitor who just taught the funnel their interests must do that work again. An interrupted setup loses its draft; the product benefit arrives late.

Change. Prefill a resumable onboarding draft from validated funnel intent. Confirm the goal and guidance in one short step. Move focus and username to the moment they are useful, with username required before posting. Ask risk questions before paper execution.

Acceptance. Complete signup, close/reopen, and resume the same step with the same choices. Keep identity and intent separate until an account exists. Do not silently promote readiness from a marketing persona.

F02 - P1 / Promise only the capabilities available on this build

Evidence. Kai plan says brokerage can be connected any time in Account. Account states paper is the only mode and a broker comes later. The training hero promises seven days although only one lesson is authored. [S2, S6, S8]

Why it matters. Beginners cannot reliably distinguish a product promise, a simulation and a future capability.


Change. Use shared capability data for the site, onboarding and app. Say Practice with paper money in the action itself. Offer only available lessons and label the broader curriculum as planned until it is usable.

Acceptance. Every advertised capability has a working destination for that account and platform. Availability changes update both the offer and the destination.


Recommended first minute

After the reveal, let the beginner complete one short visual interaction and receive useful feedback. Let the swing trader inspect one real or clearly labeled example setup. Let the experienced visitor ask Kai to mark one chart. Account setup should preserve that context.

Home: the first useful object is too low



Beginner Home proof, committed 7 September. Training follows the wake-up.



Trade Ready Home proof, committed 7 September. Mostly the same objects, reordered.

F03 - P1 / Make the opening stage-specific

Evidence. Wakeup renders before training in every stage. homeOrderFor only moves training above or below the conversation items, not above the opening briefing. The beginner-pick object is explicitly unbuilt. [S1]

Why it matters. A beginner reads market status and an explanation about an empty watchlist before reaching something they can use.

Change. Open with a one-line Kai greeting and the next action: a lesson for beginners; a company for investors; a setup or position for traders. Put the full briefing behind Read the briefing. Keep one persistent way to ask Kai.

Acceptance. At a 390-pixel phone width, the first actionable object and its button are visible without scrolling. In five seconds, members can say what to do next.

Kai: repair continuity before adding polish

F04 - P0 / Conversation selection changes the label, not the underlying thread

Evidence. Home stores the selected ConversationRow and changes seed text. useKaiWall receives only mode and seed, keeps its own convold ref, and does not receive the selected row ID. New conversation also changes seed without resetting this ref. No saved-message loader is wired into this selection. [S1, S5]

Why it matters. The screen can say it is resuming one conversation while subsequent turns use another conversation created during the mounted Home session. The visible history is also not restored.

Change. Make conversationId an explicit input to the wall. Load the selected transcript, abort an old stream before switching, and create a new server conversation for New. Pin symbol/setup context to the selected thread.

Acceptance. Create A and B with distinct topics; switch between them and send a turn. Check transcript and server ID. Refresh, resume B, then start New; no history or response may cross between them.

F05 - P2 / Recovery and control should be visible

Evidence. The shared composer clears on send and disables editing during streaming. Home appends generic stream errors; no user-facing Stop control is exposed by useKaiWall. The contextual sheet is a separate thread implementation. [S5, S13]

Why it matters. People can feel trapped waiting, and must reconstruct work after an unsuccessful request. Different Kai entry points can feel like different assistants.

Change. Use one interaction contract: context label, composing, thinking, streaming, stop, retry, and saved history. Preserve failed input; retry the same request context. Provide short suggested questions tied to the object on screen.

Acceptance. Test a failed request, interrupted stream and thread switch. Input survives, a retry is discoverable, and a late response cannot appear in the wrong context.

Trade ideas should explain themselves

F06 - P1 / Entry, stop and target are hidden behind expansion

Evidence. StandardAlertCard puts its stock level strip inside the open condition. Expanded labels are 8.5 and values 12 logical pixels. An options-led card has a visible contract object, but this does not make its information equivalent to a stock setup. [S4]

Why it matters. The user must inspect the card before learning its most useful facts. A strong visual setup may look like a generic alert headline in the list.

Change. Make SetupPreview the default object: company identity, setup type, availability state, a clear price map, entry / stop / target, a risk-reward ruler and one short explanation. Collapse evidence, score breakdown and longer discussion, not the decision essentials.

Acceptance. Without expanding, a novice can identify the proposed entry, where the idea fails, the target and whether this is only an idea or an active order. Key values remain legible at large text sizes.

Different instruments need honest variants


A complete swing setup, unusual options activity, and a long-term company thesis are different products. Keep shared identity and state styling, but do not manufacture absent grades, stops, targets, or reward ratios to fill a beautiful template. Use a visible No exit plan supplied state where relevant.

For options activity, show the contract, expiration, cost basis when supplied, and what was detected. Offer Explain this signal before offering a plan. A beginner who understands a stock setup should not be assumed to understand a contract.


A good five-second card

NVDA. Swing setup. Waiting for entry. A clear price path with Entry, Stop and Target. Risk 1R / potential reward 2.6R. One sentence explaining the setup. One primary action matched to its current state. These values are illustrative, not a current recommendation.

Charting: keep the strong structure



Look / Decide / Take proof from 4 September. The current route retains this structure.



Historical ungraded-state proof. Current source preserves an explicit No graded setup branch.

F07 - P1 / Reduce competition on the chart

The current implementation already shares SymbolChart and separates Look, Decide and Take. Preserve that. Chart heights are fixed at 352 for Look and 150 for Decide / Take. Earlier proof illustrates closely packed labels; verify their current collision behavior in the release build. [S4]

Give Look the largest usable chart area after safe areas and controls. In Decide, show a clean static setup map with only selected levels. A tap on a level should highlight one annotation and reveal one explanation. Expand to advanced drawing tools when requested.

Use Review paper trade or Build a plan as forward actions. Keep the context visible when Kai opens. Avoid using Take it as the only explanation of what the next tap will do.

Acceptance: on a small phone and large text, prices and line labels do not collide, chart gestures do not trap scrolling, and a beginner can explain one highlighted level without reading a paragraph.

Make the trade lifecycle unmistakable

F08 - P1 / Unify execution wording and visual state

Evidence. Portal confirmation uses Send it; the dedicated review screen uses Place paper order / Place paper exit. The backend already supports preview, risk blockers and submit, and the portal distinguishes accepted from filled. Preserve those foundations. [S4, S9]

Why it matters. A friendly phrase is insufficient when it changes an order. Ideas, watched alerts, drafted plans, submitted orders and positions need distinct identities.

Change. Use explicit actions: Track setup, Activate alert, Review paper order, Place paper order, View order, Review position. The state strip should show where this object is now and one next action. Carry the same symbol, direction and account mode throughout.

Acceptance. The same action has the same name in every entry point. Paper is visible at final confirmation and receipt. Double-tap, timeout and reconnect cannot create ambiguous submissions; authoritative order status resolves uncertainty.

Object state	What the member should understand
Idea / watching	No order exists. Explain whether Kai is monitoring and how delivery works.
Planned / ready	A plan is saved. Review the latest quote and risk before confirming.
Order pending / partial	The order was submitted; show filled quantity and the remaining state.
Position active	Show planned risk, current position, stop/target status and what changed.
Closed / invalidated / expired	Explain the reason. Offer a debrief or a useful replacement action.

Keep source timestamps and delayed-data indicators adjacent to quotes. Present stop-based risk as planned risk, not a guaranteed loss ceiling. Do not reuse an old illustrative setup as a current signal. [S4, S9]

Training: finish the promise first

F09 - P0 / The seven-day journey cannot currently be completed

Evidence. The curriculum defines 31 lesson nodes; only d1l1 has `hasContent: true` and only that lesson exists in `LESSON_CONTENT`. All later lessons are unavailable. The first lesson also includes a transparent video placeholder. Day 2 and Day 7 readiness gates require unavailable content. [S6]

Why it matters. A motivated beginner reaches an implementation boundary just as the app asks them to trust the learning path. No amount of progress-ring polish can solve this.

Change. For the immediate release, advertise the playable foundation lesson and provide a useful next step after it. Before promoting a complete seven-day program, author and validate every required lesson, assessment and handoff. Keep future days clearly labeled as planned.

Acceptance. A fresh account can complete every advertised required step through the final handoff. No required path ends at Soon, a disabled next lesson, or an unauthored destination.

F10 - P0 / Progress belongs to the device, not the learner


Evidence. `TrainingProvider` reads one constant `AsyncStorage` key, `ccai.training.profile.v2`, at mount. It does not use an account ID. The provider wraps the app; no training reset on sign-out was found. Readiness evidence is later sent from this local profile. [S6, S3]

Why it matters. Different accounts on the same install can inherit learning progress. A new device cannot recover that progress. This can also contaminate readiness evidence, so it is a product-integrity issue.


Change. Store learning state per authenticated learner with durable server sync. Isolate guest practice, clear in-memory state on account changes, and migrate local progress only with a verified owner. Include curriculum version and idempotent completion records.

Acceptance. Account A and B on one device stay isolated. A resumes on device two. Offline completions merge once. Clearing storage or reinstalling does not erase server-backed progress.

Make learning tactile and forgiving



The ownership grid is a useful custom mechanic; prose still precedes it.



Sorting is already interactive. Retain and refine this approach.

F11 - P1 / Save the moment, not only the finished lesson

The first authored lesson is a 14-screen sequence. LessonRunner persists its result at completion; index and answers are local component state. Leaving mid-lesson restarts that run. [S6]

Checkpoint each interaction, add Back and Resume, and distinguish practice feedback from assessed mastery. Keep explanations close to the object: tap company tiles, mark a level, move a stop, compare two plans. Offer a hint before revealing the answer.

For the ownership mechanic, always identify the denominator. Ten shares means 10% only in a company with 100 total shares. The current million-share lesson correctly avoids that shortcut; preserve its conceptual accuracy.

Acceptance: leave at screen 10, reopen and resume at screen 10 with answers intact. Finish once, reopen and see the same result. A storage failure offers recovery instead of silently presenting unsaved success.

Progress should explain achievement

F12 - P2 / Too many progress systems compete for meaning

Evidence. Training shows overall percentage, day completion, nine mastery bars and competency signals. Repeating a lesson adds masteryGain again even when its completed ID is already stored. Stages can also start at Trade Ready from self-report, while belt points depend on resolved calls and trade outcomes. [S3, S6, S7]

Why it matters. Users cannot easily tell what they learned, what remains, and whether a label reflects knowledge, experience or results. A screen of zero bars can make an early success feel small.

Change. Lead with the next demonstrated skill and a concise achievement: You can explain ownership. Put detailed metrics behind View progress. Make mastery evidence-based and repeat-safe; keep practice repetition useful without turning repeat taps into competence. Distinguish self-reported placement from course completion.

Acceptance. Repeating the same result does not inflate mastery. A member can explain why they have a readiness stage and how the next milestone is earned. A stage transition has a clear, dismissible explanation.

Keep belts aspirational, but explain what they measure

The current belt rules reward outcomes of resolved calls and Kai trades. That is not the same as a demonstrated learning competency or risk discipline. Do not present a Black Belt as proof of safety or as a guarantee that someone should be copied.

Show one belt and progress toward the next, with scoring rules and enough context about the underlying record. If discipline is meant to earn progression, it must be measured by the scoring system before the UI claims it. Avoid rewarding more trading simply to fill a ring.

Connect practice to improvement

After a paper trade closes, make the debrief easy to find from the receipt, position and Home. The existing debrief list already prioritizes trades awaiting review. Use one concrete improvement from that review to suggest the next practice exercise. [S9]

Community should feel welcoming

F13 - P1 / Reading another room changes trading mode

Evidence. Community uses ModeSegmented for Day / Swing / Invest. That control writes the global mode through PUT /mode; changing rooms therefore also changes Home and the second tab. The Beginners pill only changes the room, creating two different behaviors in one header. [S7]

Why it matters. A swing trader checking an investing conversation can return to a changed research screen and different recommendations without intending to change their goal.

Change. Make room selection local navigation. Keep global trading mode in an explicit preference control, with clear confirmation of its effect. Remember the last room independently of the trading goal.

Acceptance. Visit every room and return to Home and Alerts: primary_mode is unchanged. Explicitly change mode: Home and the second tab update once, with an understandable acknowledgement.

F14 - P2 / A dense header and empty room ask too much of newcomers

Evidence. The beginner proof squeezes title, room, member count, search, Beginners, three modes and members into the header. The empty room invites the member to say what they are watching; it also exposes calls, the board and circles. [S7]

Why it matters. A novice may have no ticker to discuss and no confidence to post. The app introduces several community concepts before showing one useful conversation.

Change. Use a readable room title and one room switcher. Show one real pinned welcome or active discussion, two optional question starters and a clear composer. Surface calls and rankings as secondary destinations. Provide a beginner-safe prompt such as Can someone explain a stop?

Acceptance. First-time visitors can find the room name, ask a question and see who replied. An empty room offers a useful start without fabricated members, activity or conversation.

Reuse the trade object in the conversation

PinnedTradePreview and ConversationPreview can connect a setup to its discussion. Preserve authorship, timestamps, replies, moderation and reporting. Keep personal position sizes private unless explicitly shared; the existing sharing defaults already support restraint.

Research and Account need clearer jobs

F15 - P2 / The investing desk still assumes market literacy

Evidence. The desk uses compact ticker rows, theme labels and states such as armed, cooled, triggered and extended. Its idea grade deliberately differs from an actionable trade grade. [S11]

Why it matters. A person who chose long-term investing may still meet trading-like vocabulary without a simple understanding of the company.

Change. Lead with company identity, what it does, a stated horizon, why it is on the desk and the main risk. Make the grade explicitly an Idea grade. Explain each watch state in plain language on demand, then offer Watch company or Ask Kai.

Acceptance. A beginner can explain what the company does and why it is being watched without opening a long thesis. Watching a company does not imply an order or a complete trade plan.

F16 - P2 / Account is a long feature directory with surprising controls

Evidence. Account contains identity, sharing, Kai profile, risk rules, many destinations, brokerage, memory, adherence, sign-out and legal controls. Trading mode and experience rows can cycle values directly; the mode chip instead opens a chooser. [S8]

Why it matters. Important tasks are difficult to find, and the same setting behaves differently depending on where it is changed.

Change. Lead with Profile, Guidance, Practice account, Notifications and Plan. Put learning progress and debriefs where their journeys naturally end as well as in Account. Use explicit selection sheets for settings with more than two values. Show save success and recover failed writes.

Acceptance. Find and change guidance, notification delivery and practice balance without scanning the full account page. A setting never changes just because the user expected its row to open details.

Preserve useful account protections

Keep memory review/deletion, clear trade-sharing controls, the paper-reset confirmation and account-deletion confirmation. Improve discoverability without weakening those flows. Do not expose a simulated development action in a release build. [S8]

Trust depends on accurate states

F17 - P1 / Entitlement presentation can disagree with the data

Evidence. Subscription filters premiumOnly from features where included is false, but displays Also open to you and green checks for non-free plans. This can label an excluded feature as available whenever a paid plan has exclusions. [S8]

Why it matters. A paying member may be told a feature is included immediately before being blocked from it.

Change. Render each entitlement from its own included value. Share this contract across tab locks, feature actions and the plan screen. Preserve capability-specific alternatives, such as the existing Ask Kai action when Trade is locked.

Acceptance. Verify free and every paid tier with at least one excluded feature. The plan list, screen gate and API agree. A service failure is not displayed as a free-plan downgrade.

F18 - P2 / Refresh attention and distinguish quiet from unavailable

Evidence. The tab badge reads alert attention once on mount. AlertsBoard has quiet-state routes and error text; Community has its own loading/error presentation. Day-trade availability is a producer switch in source, not proof that a new idea exists every session. [S1, S4, S7]

Why it matters. An outdated attention dot loses credibility. No alert, a failed request and an expired setup have different implications for a user.

Change. Refresh badges from the same acknowledged alert state used by the board. Create shared loading, quiet, stale, offline, blocked and failed patterns. Show Last checked when the service can support it; offer practice, watchlist review or a relevant discussion on quiet days.

Acceptance. Opening or acknowledging attention clears the badge correctly. A failed load cannot be mistaken for a verified empty list. Reconnect refreshes without losing the selected context.

Notification foundations are worth retaining

The current PushPrimer asks after activating an alert; DeliveryPanel distinguishes denied permission and unsupported environments. Keep that timing. Use a short delivery status beside the monitored alert, and test return-from-notification into the exact alert/position, including session expiry. [S12]

Readability is part of the product

F19 - P1 / Accessibility controls are saved but not applied app-wide

Evidence. Settings saves `accessibility.text_scale` and `reduced_motion`. Searches of the mobile source found those fields only in settings, types, fixtures and adapters. T renders its fixed size without this preference. Motion hooks in Home/Community read OS preferences instead. [S8, S13]

Why it matters. A user selecting Larger or Reduce motion can see a saved preference without getting the promised change.

Change. Use a shared accessibility provider that combines system settings with the member preference. Apply type scaling through text primitives and define motion behavior for sheets, chart annotations and feedback. Do not disable system text scaling.

Acceptance. Changing each setting visibly changes Home, a lesson, a setup and a sheet; it survives relaunch. With OS reduced motion enabled, no app setting re-enables nonessential motion.

F20 - P1 / Replace tiny, low-contrast information

`color.dim` (`#6E675F`) over the base `#0B0B0E` calculates to about 3.53:1; small dim copy therefore falls below the 4.5:1 normal-text target. Actual translucent surfaces need separate measurement. Examples include the alert mode note, metadata, stage tags and lesson footers. Important text is often 8.5-12 logical pixels. [S13]

Use 15-16 as the proposed body default, 13-14 for useful metadata, and 22-30 for key trade values. Reserve small text for genuinely secondary captions. Keep a consistent sans-serif fallback on web: the inspected public welcome/signup rendered in a serif face even though the computed family requested Space Grotesk. The font-loading cause and native behavior remain unverified.

Aim for 44 x 44 logical-pixel interactive targets as a product convention, with usable spacing and focus states. WCAG AA web target minimum is 24 CSS pixels with exceptions; enhanced guidance is 44. Test 200% text resizing, screen-reader order, chart text alternatives and keyboard focus. These are launch checks, not a claim of certification. [W1-W3]

A small custom component system

Reuse the approved visual direction across app and funnel. The components should carry a stable data contract and real behavior, not only an impressive fixture.

Component	Purpose	Where to use it
TradeMap - existing kit	One price geometry with entry, invalidation and target; responsive labels and text alternative.	Setup detail, practice, pinned trade.
RiskRewardRuler - existing kit	Explain relative planned risk and reward; show dollars only with valid size.	Setup, confirmation, lesson.
SetupPreview - existing kit	Essential setup facts and one state-driven action.	Alerts, Home, search, walkthrough.
TradeStatusStrip - existing kit	Idea / monitoring / order / position state; timestamp and reason.	Alerts, detail, notification, receipt.
KaiAnnotation - existing kit	One short contextual explanation connected to an object.	Chart, lesson, order review.
PinnedTradePreview + ConversationPreview - existing kit	Connect a real setup to a real human discussion.	Community, setup detail, funnel demo.
LessonStage - proposed	Reusable tactile task, answer feedback, hint, checkpoint and resume.	Training and beginner demo.
MilestoneRing - proposed	One demonstrated achievement and next unlock; truthful progress.	Training, Home, profile.
CapabilityState - proposed	Consistent unavailable, stale, quiet, loading and entitlement states.	Every data-driven feature.

Integration rules: preserve source timestamps; differentiate stock, option and investment data; accept missing fields without fabricating them; map lifecycle values explicitly. The marketing demo uses clearly labeled sample data through the same presentation contract.

Keep the current dark identity, warm text, cyan market data, violet Kai and volt primary actions. Reduce repeated borders and nested surfaces. Use motion to explain a state change or connect two objects, with an immediate reduced-motion equivalent.

Conversion: preserve the interactive idea

F21 - P1 / The funnel ends at an email draft

Evidence. The latest site has persona, cinematic value reveal, branching demos and a personalized recap. `getAppHref` carries path, interest and priority to `/get-the-app`, which opens mailto. No account, subscription or CRM record is created by that handoff. [S10]

Why it matters. The most engaged visitor must leave the flow and send a message manually. Their answers are included in that draft but are not automatically available in onboarding.

Change. Keep the existing conversational funnel. If the product remains private, end with a short native early-access form and confirmed request state. If signup is ready, hand off a validated intent token and resume the chosen experience after authentication. Let users change their answers.

Acceptance. A completed request produces a confirmed record or a recoverable error, including in social in-app browsers with no configured mail client. Onboarding can retrieve the submitted intent for the correct person.

Show the product the user can actually receive

Use the same trade and conversation components in the reveal, the demo and the app. The current app is paper-first; training availability and plan access must be reflected in the demo recap. A cinematic promise followed by an unavailable feature is a conversion and retention problem.

Free and paid journeys

Free members still need an achievable first success. A Trade gate currently preserves an Ask Kai alternative. Make the free promise clear before a sequence relies on a paid feature. Existing native purchase restrictions are deliberate product constraints in source; validate each platform/storefront before changing commercial navigation. This audit does not treat source comments as a definitive platform-policy ruling.

No conversion rate can be inferred from code or screenshots. The likely improvement is reduced confusion and fewer interrupted handoffs; measure it with the experiments on the next page.

Validate ease, not more trading

Test group	Task	Success evidence
Brand new	Explain a share; leave and resume the lesson.	Correct explanation, same saved step, no help needed to continue.
Investor	Understand one company and add it to a watchlist.	Can state the business, horizon and one risk; knows no order was placed.
Developing	Read a setup and build a paper plan.	Correct entry / invalidation / target; can explain planned risk.
Trade Ready	Analyze a chart, confirm a paper order and find its status.	Understands submission versus fill; finds the position and debrief.
All groups	Switch community rooms, return to work; change text size.	Goal remains unchanged; work and accessible layout are preserved.

A practical first usability round

Recruit 3-5 people per readiness group, including a long-term investor, with at least some arriving from a social link. Observe tasks without coaching. Capture completion, wrong turns, misunderstood labels and confidence before/after. This is a qualitative starting round, not a statistically conclusive conversion study.

Instrument the whole handoff

Track `persona_selected`, `value_reveal_completed`, `demo_interacted`, `interest_selected`, `early_access_submitted`, `signup_completed`, `first_value_completed`, `lesson_resumed`, `setup_understood`, `alert_activated`, `paper_order_confirmed` and `debrief_viewed`. Include build, stage, goal, entry source and relevant state. Avoid collecting chat text or account values just to measure funnel steps.

Use conversion denominators explicitly: early-access submission / eligible recap viewers; first meaningful success / onboarded accounts; lesson resumes completed / interrupted lessons resumed. Segment free versus paid and social in-app browser versus full browser. Do not optimize number of trades as a UX success metric.

Proposed targets to test, not measured results

Users can identify the next action within five seconds on a primary screen. First useful interaction is reachable within a minute after identity is established. No advertised required path ends in unavailable content. No wrong-account progress or wrong-thread reply is acceptable.

Build order and release gates

Batch	Work	Exit criteria
1 / Repair trust	F04 conversation identity; F09 curriculum promise; F10 learner ownership; F19 accessibility settings; F17 entitlements.	Thread IDs, account state and entitlements match visible claims. All advertised lessons are completable. Accessibility controls work.
2 / First useful experience	F01-F03 onboarding and Home; F06 trade preview; F08 explicit execution language; F11 checkpointing.	Each persona reaches one understandable action; beginner can resume; core setup values are visible without expansion.
3 / Connect the app	F13 room independence; F14 discussion; F15 desk; F16 Account; F18 shared state and badge refresh.	Moving between features preserves goal and context; quiet, failed and blocked states are distinct and useful.
4 / Finish the experience	F05 recovery; F07 chart refinement; F12 milestones; F20 visual/accessibility refinement; F21 signup handoff.	Real-device usability round passes; funnel and app share components and capability rules.

Use the component-kit integration inside batches 2-4, starting with SetupPreview and TradeMap. Do not combine a full visual rebuild with changes to order semantics in one unreviewable release. Each batch should produce a complete user journey that can be demonstrated.

Required validation before a release claim

Signed-in iOS and Android release builds; desktop and mobile web; social in-app browser; small phone and enlarged text; VoiceOver / TalkBack; keyboard and sheet focus; slow or dropped connection; market closed / delayed quote / missing setup; free and paid accounts; account switch and two-device progress; interrupted orders and restored sessions.

Run these against the real authentication gate and API. The source notes prior cases where fixture previews skipped routing failures, so fixture screenshots alone cannot close this gate. [S13]

The first implementation ticket

Fix Home conversation identity and learner-owned persistence, then reconcile advertised training availability. In parallel product design work, finalize one accessible SetupPreview using the existing kit. That gives the redesign a reliable foundation and a reusable visual standard.

Evidence index 1 / 3

S1 / Home and navigation

[mobile/app/\(tabs\)/home.tsx](#)

Wakeup ordering, selection and seed wiring

[mobile/features/stage/home-order.ts](#)

Training ordering only

[mobile/app/\(tabs\)/_layout.tsx](#)

Badge fetch and tab entitlements

S2 / Onboarding

[mobile/app/\(onboarding\)/start.tsx](#)

Placement and goal

[mobile/app/\(onboarding\)/personalize.tsx](#)

Second experience question

[mobile/app/\(onboarding\)/kai-plan.tsx](#)

Capability promises

[mobile/lib/session.tsx](#)

Default draft and lifetime

S3 / Readiness

[api/lib/stage/rules.ts](#)

Self-report placement and promotion

[mobile/features/stage/labels.ts](#)

Stage meaning

[mobile/features/stage/useStageEvolution.ts](#)

Local evidence reporting

S4 / Alerts and current Trade

[mobile/features/alerts/AlertCard.tsx](#)

Collapsed level strip

[mobile/features/alerts/AlertsBoard.tsx](#)

Quiet and error states

[mobile/app/trade/\[symbol\].tsx](#)

V2 is current route

[mobile/features/portal2/TradePortalV2.tsx](#)

Chart and interaction layout

[mobile/features/portal2/Take.tsx](#)

Confirmation and receipt

S5 / Kai conversations

[mobile/lib/useKai.ts](#)

Home convold and separate sheet thread

[mobile/features/home/useConversations.ts](#)

Drawer list and metadata

[mobile/lib/api.ts](#)

Conversation API methods

Evidence index 2 / 3

S6 / Training

[mobile/features/training/curriculum.ts](#)

31 nodes, 1 authored lesson

[mobile/features/training/store.tsx](#)

Global storage key and repeated mastery gain

[mobile/features/training/engine/LessonRunner.tsx](#)

14-screen run and completion persistence

[mobile/features/training/content/day1-lesson1.ts](#)

Ownership mechanic and video placeholder

[mobile/app/training/index.tsx](#)

Seven-day promise

[mobile/features/training/gates.ts](#)

Required day gates

S7 / Community and belts

[mobile/app/\(tabs\)/community.tsx](#)

Room switch uses global mode

[mobile/features/social/belts.ts](#)

Belt presentation

[api/lib/social/belts.ts](#)

Outcome-based scoring explanation

S8 / Account and access

[mobile/app/\(tabs\)/account.tsx](#)

Long account surface and setting controls

[mobile/app/account/settings.tsx](#)

Saved accessibility and guidance preferences

[mobile/app/account/subscription.tsx](#)

Excluded features labeled available on paid

[mobile/app/account/paper.tsx](#)

Practice reset

[mobile/app/account/delete.tsx](#)

Deletion confirmation

[mobile/features/portal2/TradeLocked.tsx](#)

Existing locked-state alternative

S9 / Orders and improvement

[mobile/app/order/review.tsx](#)

Explicit paper order review

[mobile/features/portal2/useTake.ts](#)

Preview / submit / receipt flow

[api/lib/execution/preview.ts](#)

Risk blocker checks

[mobile/app/position/index.tsx](#)

Position navigation

[mobile/app/debrief/index.tsx](#)

Post-trade learning entry

S10 / Funnel and conversion

[site/components/Walkthrough.tsx](#)

Branching demo and recap

[site/sim/handoff.ts](#)

Intent query parameters

[site/app/get-the-app/page.tsx](#)

Mailto conversion boundary

Evidence index 3 / 3

S11 / Investing

[mobile/features/desk/Watchlist.tsx](#)

Watchlist presentation

[mobile/features/desk/ui.tsx](#)

Idea grade meaning

S12 / Notification delivery

[mobile/features/notifications/PushPrimer.tsx](#)

Contextual permission primer

[mobile/features/notifications/DeliveryPanel.tsx](#)

Delivery and blocker states

S13 / Shared UI and release gate

[mobile/ui/tokens.ts](#)

Colors and screen offsets

[mobile/ui/Text.tsx](#)

Fixed text sizing

[mobile/ui/fonts.ts](#)

Font loading and family names

[mobile/ui/Composer.tsx](#)

Send and editing behavior

[mobile/ui/Sheet.tsx](#)

Sheet interaction

[mobile/app/_layout.tsx](#)

Session gate and provider placement

Live observation

Welcome and signup inspected at <https://cheatcode-ai-app.vercel.app> on 8 September 2026. No authentication was attempted. The public page lacked the expected sans-serif visual appearance and stretched phone-style forms across desktop width. This is a deployment observation, not proof of the cause or of native behavior.

Accessibility references

W1 / [W3C contrast minimum](#)

W2 / [W3C target size minimum](#)

W3 / [W3C target size enhanced](#)

Design proposals, severity and sequencing are audit judgments. Screenshots are repository proof artifacts, not fresh authenticated captures. No conversion statistics or performance results were available.